

Verifying Distributed Algorithms with Interactive Theorem Provers: A Survey of Approaches

MAXWELL MOIR, University of Canterbury, New Zealand

Distributed algorithms are a fundamental aspect of modern computer science. They are notoriously difficult to correctly design and implement, and must conform precisely to their specification. Formally verifying these algorithms guarantees their correctness, although this is often a tedious and time-consuming process. Despite this, proving the correctness of a program is not only theoretically satisfying but is often necessary in mission-critical applications. This paper provides a survey of existing literature on the formal verification of distributed algorithms using interactive theorem proving. It provides taxonomies of common tools, proof techniques, and algorithms. Finally, it identifies common challenges in the field and potential research directions that might improve the feasibility of verifying these algorithms in industry.

CCS Concepts: • **Theory of computation** → **Logic and verification**.

Additional Key Words and Phrases: distributed systems, formal verification, proof assistants

ACM Reference Format:

Maxwell Moir. 2026. Verifying Distributed Algorithms with Interactive Theorem Provers: A Survey of Approaches. *J. ACM* 000, 0, Article 000 (2026), 7 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

In an increasingly interconnected world, distributed algorithms play an integral role in much of the software we use daily. Ensuring these algorithms function as intended protects against failures that may be technically, economically, or even medically critical.

1.1 Motivation

Distributed systems are famously difficult to implement correctly. Their divided and concurrent nature can often lead to subtle bugs that are difficult to reproduce. Additionally, errors can propagate throughout a system, and partial failures can lead to unexpected negative outcomes. For this reason, it is often unrealistic to achieve total coverage with traditional testing. Formal verification can provide a strong guarantee of correctness in these situations. Interactive theorem proving as a verification method is especially strong for distributed systems, as it can verify correctness without explicitly exploring states. Theorem proving also allows for flexibility in the parameterisation of systems, such as verifying an algorithm for any arbitrary number of nodes, N . For this reason, this survey will explore the verification of distributed algorithms with interactive theorem provers.

1.2 Contributions

This survey aims to make the following contributions. It will provide a taxonomy of theorem proving approaches and frameworks for formal verification. It will discuss several common distributed

Author's Contact Information: Maxwell Moir, University of Canterbury, Christchurch, New Zealand, mmo199@uclive.ac.nz.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM 1557-735X/2026/0-ART000

<https://doi.org/XXXXXXX.XXXXXXX>

algorithms and efforts that have been made to verify them. Finally, it will give a comparative analysis of these major tools, frameworks and methods.

2 Methodology

2.1 Literature Search Strategy

To identify relevant literature, this survey employed a systematic literature search strategy. Two major academic databases were selected: the ACM Digital Library and IEEE Xplore. These databases were chosen for their broad coverage of research in formal verification and distributed systems. They were accessed on April 17, 2026.

To ensure the quality of the review, the PRISMA framework was used to inform the search strategy. The strategy consisted of constructing a query followed by database searching and screening. First, relevant keywords were curated based on the topic of this survey, and a search query was constructed. The query used is as follows.

("proof assistant" OR "theorem prover" OR "mechanized verification" OR "formal verification") AND ("distributed protocol" OR "distributed algorithm")

Selected papers were limited to peer-reviewed English-language publications. There was no strict lower bound imposed on the publication date, although results were naturally concentrated in the last 20 years. This query yielded 213 results from the ACM Digital Library and 49 results from IEEE Xplore, which were exported for further screening.

2.2 Inclusion and Exclusion Criteria

These results were then screened by title and abstract with standard inclusion and exclusion requirements. Included papers focused on the verification of distributed algorithms using theorem proving. Duplicate papers within and between databases were removed. Additionally, papers were excluded if model checking or manual proof were the primary verification technique. This reduced the number of results to 37 for ACM and 15 for IEEE Xplore.

2.3 Citation Chaining

Finally, the reference lists of the selected papers were scanned, and fundamental or commonly cited papers were added. Google Scholar was used as a supplementary tool for citation chaining. Any additional papers identified this way were included only if they met the above criteria or were primarily concerned with a related topic to be explained. After screening and citation chaining, 17 papers were included in the final survey.

3 Preliminary

3.1 Distributed Algorithms

Distributed computing is a fundamental field in Computer Science. It concerns communication among distinct autonomous components in a system. These patterns often arise when there is a network of multiple devices communicating over unreliable networks.

There are many challenges in distributed computing that can be the targets of formal verification. Several of these problems will be elaborated upon below, such as consensus, fault tolerance, and mutual exclusion protocols.

3.2 Interactive Theorem Provers

Interactive Theorem Provers, or proof assistants, have seen a recent surge in popularity among computer scientists and mathematicians alike. Proof assistants are tools that assist the human development of formal proofs. These formal proofs can then be automatically checked by the

machine, ensuring correctness at each stage of the process. Theorem provers are extremely useful in the formal verification of protocols, where they can be used to ensure specific properties of a system hold. For this reason, proof assistants such as Coq, Lean, Isabelle, and Dafny have become popular for verifying distributed algorithms.

3.3 Verification

Verifying the correctness of a program requires proving that specific conditions about a system hold in every state. The two most notable property categories of a system are “Safety” and “Liveness”. A safety property ensures protection against unfavourable events, whereas liveness properties make sure that the program functions as intended. Although these are both important properties, liveness is significantly more difficult to prove than safety and typically require more sophisticated techniques such as temporal reasoning.

Fault tolerance is another notable aspect of ensuring a program works as intended. Byzantine fault tolerance, specifically, is a common problem in distributed computing.

4 Approaches to Verifying Distributed Algorithms

4.1 Tools and Common frameworks

There are many languages and frameworks used to verify distributed systems. The most commonly used tool in the literature is Coq (recently renamed Rocq). There are multiple frameworks implemented for Coq that specifically focus on distributed systems, such as Verdi [Wilcox et al. 2015] and Disel [Sergey et al. 2017]. These frameworks both use Coq, but have different focuses. As described in the Disel specification, Verdi focuses on vertical composition, meaning it aims to prove correctness across various fault models. Disel, on the other hand, is primarily used to establish safety invariants between separate interacting protocols.

Dafny is another verification-aware programming language that allows the verification of distributed algorithms. Iron Fleet [Hawblitzel et al. 2015] is a framework built on Dafny that is designed for this specific purpose. Another regularly used theorem prover is the TLA+ Proof System, or TLAPS [Cousineau et al. 2012]. This was created by Lamport as a supplementary tool to his TLA+ formal specification language. TLAPS can be used to prove correctness conditions about systems that have been formalised in TLA+. A more contemporary tool is Lean4. In 2025, the Veil [Pirlea et al. 2025] framework was introduced as a basis to verify distributed protocols specifically. Due to its novelty, there have been relatively few published attempts to verify algorithms with Lean4 outside of the Veil specification itself. Veil has both automated and interactive theorem proving capabilities built in. This means a user can interactively prove properties using native Lean tactics or use Veil’s built-in SMT solvers.

These languages and frameworks all have advantages and disadvantages. IronFleet produces executable C# code and has a stronger focus on verifying liveness than the other frameworks. Lean4 with Veil is lightweight and expressive, but Veil is still a research prototype, unlike Dafny and Coq, which have been tried and tested in industry. Another aspect where these tools differ is their level of automation. IronFleet, TLAPS, and Veil all have SMT solver support, meaning they can use tools such as Z3. The Coq-based tools typically rely more on interactive proof, with less built-in automation.

Overall, these tools each have their own intended use-case and reflect trade-offs in ease of development, expressiveness, automation, and industrial maturity.

4.2 Proof techniques

There are several notable proof techniques used for proving correctness properties in distributed algorithms. The most common technique used with interactive theorem provers is *inductive invariants*. When using this approach, the verifier proves that a property holds in the initial state and that it is preserved by every transition. Inductive invariants can be seen used in a verification of a consensus protocol by Chand et al. [Chand et al. 2016], in which Type, Acceptor, and Message invariants are shown to hold inductively under all transitions. These invariants are then shown to prove safety, as defined in the paper.

Refinement-based proof approaches aim to bridge the gap between abstract models and executable code. Refinement mappings [Abadi and Lamport 1991] are used to create a relationship between states in a low-level system and a high-level specification, and a verifier can prove that there are corresponding transitions in each model. Correctness properties can be proven to hold in the higher-level system, ensuring that the protocol works as intended. In the later analysis of fault tolerance protocols, an example of a refinement proof will be provided.

Compositional verification reasons about a system by decomposing it into independently verifiable parts. As mentioned previously, this technique is used by both Verdi and Disel, but in different ways. In a 2024 paper, Zhao et al. [Zhao et al. 2024] introduced BYTHOS, a Coq framework for compositional verification. They use BYTHOS to verify safety and liveness properties of three Byzantine fault tolerance protocols: Reliable Broadcast, Provable Broadcast, and Accountable Byzantine Confirmer. There have been mixed opinions about the efficacy of composition as a proof technique. In his paper “Composition: A Way To Make Proofs Harder” [Lamport 1998], Lamport argues that mathematical decomposition is a more powerful approach than forming a specification as the parallel composition of separate components. The Disel specification [Sergey et al. 2017] addresses this viewpoint, arguing that the strengths of mathematical models and program logics can be utilised together.

These techniques are best understood in practice when applied to well-known distributed computing problems.

5 Case Studies

5.1 Consensus Algorithms

Consensus algorithms are used to agree on one result in a group of multiple participants. They typically assume that participants can crash, but will not act maliciously. Two well-known solutions to the consensus problem are the Paxos and Raft protocols. The Paxos protocol, originally specified by Lamport [Digital Equipment Corporation and Lamport 2019], is one method to facilitate consensus.

Multi-Paxos, a version of the Paxos algorithm in which distributed processes agree on a sequence of values, was specified and verified in a 2019 paper by Chand et al. [Chand et al. 2016]. This paper formally specified the algorithm in TLA+ and provided a proof written in TLAPS.

Consensus is also an important problem in application. For example, in blockchain systems, consensus is required between distributed, and potentially faulty, machines to ensure users can agree on the state of the system. In their 2018 paper [Pirlea and Sergey 2018], Pirlea and Sergey present a blockchain-based distributed consensus protocol, verified in the Coq proof assistant. This is an example of a modern application of a formally verified distributed protocol, in which correctness is absolutely critical for user trust in the system.

5.2 Byzantine Fault Tolerance

Byzantine fault tolerance (BFT) protocols, also introduced by Lamport et al. [Lamport et al. 2019], aim to achieve consensus in a group of nodes containing potentially malicious members. Consensus

protocols, such as Paxos and Raft, assume a crash-resistant model, but cannot handle unspecified node behaviour. Guaranteeing BFT is extremely important in decentralised systems, such as those used for data sharing or financial applications, which are often the targets of adversarial parties.

In 2022, Zhao et al. published a paper [Zhao et al. 2024] in which multiple BFT protocols were mechanically proven using compositional verification. This compositional approach, as described earlier, allowed the researchers to verify the safety and liveness properties in the protocols.

5.3 Mutual Exclusion

The mutual exclusion problem is another classic challenge in distributed computing. It was originally posed by Dijkstra in a 1965 article [Dijkstra 1965]. For a protocol to satisfy mutual exclusion, no two independent processes may be executing their so called “critical portion” at the same time.

An example of the mutual exclusion property formally verified can be found in the Verdi specification [Wilcox et al. 2015]. Here, a simplified lock service application is implemented in Verdi, with mutual exclusion verified as one of the correctness properties. The mutual exclusion property is expressed as a predicate and is shown to hold under all traces produced under a given semantics.

6 Comparison of Verification Strategies

There are three major verification paradigms that are often used for distributed algorithm verification. The other two predominant methods of formally verifying distributed systems are model checking and manual proofs. Each of these other verification approaches has its own advantages and disadvantages in relation to theorem proving.

Many solutions to the previously mentioned problems are based on manual proofs. For example, De Prisco et al. [De Prisco et al. 1997] developed an I/O automaton model and verified the Paxos algorithm. The paper contains a manual correctness proof, in which the validity of each component is proved with invariants. Manual proofs are more expressive than mechanised proofs, but have a weaker correctness guarantee because they depend on mathematical arguments formulated by humans. This means they are more prone to errors in their development or implementation, and are more difficult to maintain when scaling.

Model checking is another approach for verifying the correctness of a distributed system. Model checking expands the state space, exhaustively exploring every possible state and ensuring that the correctness properties hold in each one. This was exemplified by Tsuchiya and Springer [Tsuchiya and Schiper 2008], who used bounded model checking to verify two consensus algorithms, including Hybrid-1, an improved version of the Fast Paxos algorithm. They were able to mechanically verify agreement up to 8 processes and termination up to 11 processes. A major limitation of model checking is the state space explosion problem, which restricts it to bounded or specifically abstracted models. This often reduces the maximum number of processes in a verified protocol.

Theorem proving mitigates this issue by reasoning about the system mathematically rather than expanding states, allowing correctness conditions to be proven for arbitrary system sizes.

7 Challenges and Future Directions

One of the common criticisms of theorem-based formal verification is that it is too time-consuming to apply practically. In addition to this, verification requires specialised engineers, often with a mathematical background. Despite this difficulty, there has been industrial adoption of formal verification of distributed systems. For example, Brooker and Desai [Brooker and Desai 2025] outlined how Amazon Web Services use formal and semi-formal methods to maintain availability and security. In this article, they detail the use of formal mathematical tools such as Dafny to prove security properties in their access control language Cedar. Despite the successes at Amazon, the

authors reiterate the challenges with the steep learning curve of verification. The specialisation required for engineers maintaining formally verified systems is an ongoing challenge.

One potential direction of future research is the application of machine learning to assist with formalisation. In a 2023 paper [Liu et al. 2023], Liu et al. propose a technique based on reinforcement learning for predicting contract invariants that are useful for proving arithmetic safety. Although this paper doesn't focus on a distributed protocol, it highlights the potential that ML may hold in reducing the verification time of complex systems.

8 Conclusion

This survey provided an overview of the use of interactive theorem provers for the formal verification of distributed systems. It covered common languages, frameworks, proof techniques, and algorithms in the area and discusses their relative advantages and disadvantages. Throughout the paper, several patterns have repeatedly manifested. For example, the trade-offs between development time and the degree of correctness guaranteed. Case studies in consensus agreement and Byzantine fault tolerance have demonstrated the feasibility and importance of protocol verification.

Due to common intricacies in distributed algorithms, correctness guarantees can be especially valuable. In relation to manual proofs and model checking, interactive theorem provers act as a powerful tool in reasoning about the correctness of these systems. Finally, this survey has provided an evaluation of verification strategies, identified challenges in the area, and outlined the possible direction of future research.

References

- Martín Abadi and Leslie Lamport. 1991. The Existence of Refinement Mappings. *Theor. Comput. Sci.* 82, 2 (May 1991), 253–284. doi:10.1016/0304-3975(91)90224-P
- Marc Brooker and Ankush Desai. 2025. Systems Correctness Practices at Amazon Web Services. *Commun. ACM* 68, 6 (June 2025), 38–42. doi:10.1145/3729175
- Saksham Chand, Yanhong A. Liu, and Scott D. Stoller. 2016. Formal Verification of Multi-Paxos for Distributed Consensus. In *FM 2016: Formal Methods*, John Fitzgerald, Constance Heitmeyer, Stefania Gnesi, and Anna Philippou (Eds.). Vol. 9995. Springer International Publishing, Cham, 119–136. doi:10.1007/978-3-319-48989-6_8
- Denis Cousineau, Damien Doligez, Leslie Lamport, Stephan Merz, Daniel Ricketts, and Hernan Vanzetto. 2012. TLA+ Proofs. *Proceedings of the 18th International Symposium on Formal Methods (FM 2012)*, Dimitra Giannakopoulou and Dominique Mery, editors. Springer-Verlag Lecture Notes in Computer Science 7436 (Jan. 2012), 147–154.
- Roberto De Prisco, Butler Lampson, and Nancy Lynch. 1997. Revisiting the Paxos Algorithm. In *Distributed Algorithms*, Gerhard Goos, Juris Hartmanis, Jan Van Leeuwen, Marios Mavronicolas, and Philippas Tsigas (Eds.). Vol. 1320. Springer Berlin Heidelberg, Berlin, Heidelberg, 111–125. doi:10.1007/BFb0030679
- Digital Equipment Corporation and Leslie Lamport. 2019. *The Part-Time Parliament*. Association for Computing Machinery. doi:10.1145/3335772.3335939
- E. W. Dijkstra. 1965. Solution of a Problem in Concurrent Programming Control. *Commun. ACM* 8, 9 (Sept. 1965), 569. doi:10.1145/365559.365617
- Chris Hawblitzel, Jon Howell, Manos Kapritsos, Jacob R. Lorch, Bryan Parno, Michael L. Roberts, Srinath Setty, and Brian Zill. 2015. IronFleet: Proving Practical Distributed Systems Correct. In *Proceedings of the 25th Symposium on Operating Systems Principles (SOSP '15)*. Association for Computing Machinery, New York, NY, USA, 1–17. doi:10.1145/2815400.2815428
- Leslie Lamport. 1998. Composition: A Way to Make Proofs Harder. In *Compositionality: The Significant Difference*, Gerhard Goos, Juris Hartmanis, Jan Van Leeuwen, Willem-Paul De Roever, Hans Langmaack, and Amir Pnueli (Eds.). Vol. 1536. Springer Berlin Heidelberg, Berlin, Heidelberg, 402–423. doi:10.1007/3-540-49213-5_15
- Leslie Lamport, Robert Shostak, and Marshall Pease. 2019. The Byzantine Generals Problem. In *Concurrency: The Works of Leslie Lamport*. Association for Computing Machinery, New York, NY, USA, 203–226.
- Junrui Liu, Yanju Chen, Bryan Tan, Isil Dillig, and Yu Feng. 2023. Learning Contract Invariants Using Reinforcement Learning. In *Proceedings of the 37th IEEE/ACM International Conference on Automated Software Engineering (ASE '22)*. Association for Computing Machinery, New York, NY, USA, 1–11. doi:10.1145/3551349.3556962
- George Pirlea, Vladimir Gladshtein, Elad Kinsbruner, Qiyuan Zhao, and Ilya Sergey. 2025. Veil: A Framework for Automated and Interactive Verification of Transition Systems. In *Computer Aided Verification: 37th International Conference, CAV*

- 2025, Zagreb, Croatia, July 23-25, 2025, *Proceedings, Part III*. Springer-Verlag, Berlin, Heidelberg, 26–41. doi:10.1007/978-3-031-98682-6_2
- George Pirlea and Ilya Sergey. 2018. Mechanising Blockchain Consensus. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs*. ACM, Los Angeles CA USA, 78–90. doi:10.1145/3167086
- Ilya Sergey, James R. Wilcox, and Zachary Tatlock. 2017. Programming and Proving with Distributed Protocols. *Proc. ACM Program. Lang.* 2, POPL (Dec. 2017), 28:1–28:30. doi:10.1145/3158116
- Tatsuhiko Tsuchiya and André Schiper. 2008. Using Bounded Model Checking to Verify Consensus Algorithms. In *Proceedings of the 22nd International Symposium on Distributed Computing (DISC '08)*. Springer-Verlag, Berlin, Heidelberg, 466–480. doi:10.1007/978-3-540-87779-0_32
- James R. Wilcox, Doug Woos, Pavel Panchekha, Zachary Tatlock, Xi Wang, Michael D. Ernst, and Thomas Anderson. 2015. Verdi: A Framework for Implementing and Formally Verifying Distributed Systems. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '15)*. Association for Computing Machinery, New York, NY, USA, 357–368. doi:10.1145/2737924.2737958
- Qiyuan Zhao, George Pirlea, Karolina Grzeszkiewicz, Seth Gilbert, and Ilya Sergey. 2024. Compositional Verification of Composite Byzantine Protocols. In *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security (CCS '24)*. Association for Computing Machinery, New York, NY, USA, 34–48. doi:10.1145/3658644.3690355